

Technical concepts

Last amended: 8th July, 2026
My folder: C:\Users\ashok\OneDrive\Documents\llm
GitHub: https://github.com/harnalashok/LLMs/upload/main/install_ai_tools/misc

Contents

Technical concepts	1
What is a GPU	2
What is CUDA:	2
A basic python crash course.....	2
What are embeddings in AI	2
Made up examples	3
What are tokens in GPT or Claude?	4
What is an API?.....	4
What is a webhook?	5
API vs Webhook	5
Ollama parameters.....	5
Custom Ollama Models	8
Loop Engineering.....	9

What is a GPU

A GPU, or graphics processing unit, is a specialized processor designed to handle complex visual and mathematical tasks—especially the kind that happen all at once. A GPU might have thousands of smaller cores. These are designed to execute the same instruction (say add two numbers) across many pieces of data at once—a model known as *SIMD* (single instruction, multiple data).

In a GPU, a core is a specialised tiny unit designed to perform mathematical calculations. While a standard CPU usually has 4 to 24 highly complex cores to handle various complex system tasks, a GPU has thousands of simpler cores optimized to perform millions of *identical* calculations at exact the same time. GPU excels at massively parallel processing.

A GPU, on the other hand, might have thousands of smaller cores. These are designed to execute the same instruction across many pieces of data at once—a model known as *SIMD* (single instruction, multiple data). This is what makes GPUs so powerful for tasks like:

- Rendering every pixel in a video frame—say change gradient of every pixel by a percentage
- Running matrix multiplications in deep learning—first multiply many pairs of numbers at the same time then add some groups of numbers.

What is CUDA:

CUDA is a parallel computing platform created by NVIDIA that allows developers to speed up computationally heavy applications by harnessing the power of Graphics Processing Units (GPUs). Traditionally, computer programs run sequentially on a Central Processing Unit (CPU). CUDA allows developers to offload compute-intensive tasks to NVIDIA GPUs, which contain thousands of processing cores designed to execute millions of similar mathematical calculations simultaneously. CUDA is a software layer providing direct access to the GPU's virtual instruction set and parallel computational elements.

A basic python crash course

[Here is a basic python crash course](#), students may like to go through as refresher.



What are embeddings in AI

Embeddings are a list of numbers (called vectors) that are numerical representation of words. Think of embeddings as **GPS coordinates for abstract concepts**. On a normal map, nearby coordinates mean two places are physically close. In an embedding space, nearby coordinates mean two concepts (such as *apple* and *fruit* are *ideally* close).




Figure 1: Vectors for (cat, kitten) would be very different from vectors for (adult, child).

Think of embeddings as **GPS coordinates for abstract concepts**. On a normal map, nearby coordinates mean two places are physically close. In an embedding space, nearby coordinates mean two concepts are *ideally* close.

Made up examples

When data passes through an embedding model, it is converted into an array of numbers or **vector** containing hundreds or thousands of decimal numbers (dimensions). Each decimal number in the vector (also called dimension) tracks an abstract feature or relationship learned by the AI during training. Here is a made-up example of vectors for three words:

Vector for King:	[0.99, 0.95, 0.7]	
Vector for Queen:	[0.99, 0.05, 0.7]	
Vector for Apple:	[0.0, 0.15, 0.7]	(0.15 being preferred by males; 0.7 for being Ripe apple as against green apple)

The numbers in each of the above three vectors may track *Royalty*, *Masculinity* and *Age* respectively. Thus:

Word	Dimension 1: Royalty	Dimension 2: Masculinity	Dimension 3: Age
King	High (0.99)	High (0.95)	Adult (0.7)
Queen	High (0.99)	Low (0.05)	Adult (0.7)
Apple	None (0.00)	Neutral (0.50)	N/A (0.0)

Figure 2: Each number in the array of numbers (vector) indicates some value assigned to some feature of the word.

Here is another made up example on dimensions of *Violence*, *Entertainment value* and *Nationalistic sentiment*:

Dimensions→	Violence	Entertainment value	Nationalistic sentiment
Dhurandhar	0.9	0.7	0.9
3 Idiots	0.0	1.0	0.5
Dangal	0.1	0.8	0.95

Embedding model vectors may contain hundreds or thousands of decimal numbers (dimensions). The larger the vector space, more the semantic clarity between words but would require more computational resources to process data. Smaller the vector space, less the computational power required but then less also would be semantic similarity.

Sentence Embedding

Sentence embedding is the process of converting an entire sentence into a single numerical vector that captures its overall meaning and context.

Unlike older methods that look at sentences word-by-word, sentence embeddings analyse the structure, order, and intent of the full phrase. This allows computers to compare different sentences based on their actual ideas rather than the exact words used.

If you embed words individually, the sentences "The dog bit the man" and "The man bit the dog" look identical because they contain the exact same words.

Sentence embedding models solve this by looking at the phrase holistically. It analyses the structure, order, and intent of the full phrase. They recognize that:

- **Different words, same meaning:** "How is the weather today?" and "Is it raining outside right now?" will have vectors placed very close together in mathematical space.
- **Same words, different meaning:** "The bank of the river" and "The bank changed my credit limit" will be placed far apart because the context of "bank" is entirely different.

What are tokens in GPT or Claude?

Tokens are the basic units that commercial models use for pricing. Roughly speaking, 1 token $\approx$ 4 characters of English text, or about $\frac{3}{4}$ of a word. So "100 tokens" is approximately 75 words. Simple punctuation is usually its own token. Examples:

- "Hello!" $\rightarrow$ ~2 tokens
- "The quick brown fox" $\rightarrow$ ~4 tokens
- A typical paragraph $\rightarrow$ ~100–150 tokens

Costs are based on two components:

- **Input tokens** — everything you send to Claude/GPT: your message, any system prompt, conversation history, and attached documents.
- **Output tokens** — everything model generates in response.

Output tokens are typically more expensive than input tokens.

What counts toward your token usage?

- Your typed message
- The full conversation history (in multi-turn chats)
- System prompts (set by the operator/platform)
- Uploaded files or documents (their text content)
- Tool results (e.g., web search results passed back to the model)
- Longer conversations accumulate more input tokens over time, since prior messages are re-sent each turn.

Here's the key insight:

Unit	Approximate size
1 token	~4 characters
1 average word	~5 characters (including the space after it)
So 1 word	~1.3 tokens
Therefore 100 tokens	~75–80 words

What is an API?

An **API**, or **Application Programming Interface**, is a **software intermediary that allows two different applications to interact and share data with one another**. Think of it as a digital bridge or a translator that lets separate programs "speak" the same language without needing to know how each other's internal code is written. [Wikipedia +3](#)

The Restaurant Analogy

To understand how an API works, imagine you are a customer dining at a restaurant: [YouTube · Postman +1](#)

- **The Customer (The Client):** You sit at the table and want to order food. In technology, this is an app or website.
- **The Kitchen (The Server):** The kitchen prepares the food. This represents the backend database or system holding the data.
- **The Waiter (The API):** You don't walk into the kitchen to cook. Instead, you read the menu and give your order to the waiter. The waiter carries your request to the kitchen, retrieves the cooked food, and delivers it back to you. [GitHub +3](#)

How Web APIs Work

Most modern web communication relies on a structured **Request-and-Response** loop: [YouTube · Exponent +1](#)

1. **The Call:** Your application sends an **API request** over the internet using an HTTP protocol. [GitHub +1](#)
2. **The Endpoint:** The request travels to a specific URL path, known as an **endpoint**, which dictates what data is being asked for. [YouTube · CodeWithChris +1](#)
3. **The Method:** The request uses strict commands such as **GET** (to fetch data) or **POST** (to submit new data). [GitHub +1](#)
4. **The Security Check:** The system verifies the app's identity using a digital password called an **API key** to ensure the request is authorized. [YouTube · CodeWithChris +1](#)
5. **The Payload:** The server processes the task and sends back the data, typically formatted in a clean, machine-readable text format called **JSON**. [YouTube · Exponent +1](#)

What is a webhook?

The Doorbell Analogy

Think of a **traditional API** like a customer repeatedly calling a restaurant to ask, "Is my takeout order ready yet?". A **webhook** is like leaving your phone number with the restaurant so *they* can automatically text you the moment your food is ready.

Here, replace the 'customer' with a *web-server* (say, *App-A*), 'restaurant' with another *web-server* (say *App-B*) and 'calling' with making an API request to 'App-B'; your 'phone number' with your 'URL or your API'. Thus, instead of 'App-A' sending an API request to 'App-B' for some information, 'App-B' makes an API request to 'App-A' informing him of an event.

Polling process vs. webhook




Figure 3: Instead of CRM, repeatedly sending a message to know if the customer order is dispatched, Warehouse, informs CRM as soon as customer order is dispatched.

API vs Webhook

APIs are pull-based: Your application (App-A) actively asks for data on-demand (you request it).

Webhooks are push-based: The source application (App-B) automatically sends the data to your application the moment something occurs

Webhook is an "Event-based HTTP callback trigger"—a long name. Name, "Webhook" is punchy, precise, and easily turned into a verb (e.g., "*We can just webhook that data over*"). Webhook uses using standard web protocols (HTTP): POST/GET.

Ollama parameters

Ollama parameters are customizable settings that control how a large language model (LLM) generates text, balancing its creativity, predictability, memory, and formatting. They can be adjusted in the terminal or permanently saved in a **Modelfile**.

Key Generation Parameters

These parameters control the AI's "thinking" process and output style. 

- **temperature** (Default: **0.8**): Controls the creativity and randomness of the output.
 - **Low (e.g., 0.1 - 0.3)**: Makes the output more factual, logical, and predictable. Best for coding, math, or data extraction.
 - **High (e.g., 0.8 - 1.0)**: Makes the output more creative, diverse, and unpredictable. Best for brainstorming or writing stories.  Medium · Laurent Kubaski +4
- **top_p** (Default: **0.9**): An alternative way to control randomness, often called "nucleus sampling". It tells the model to only consider the pool of words whose combined probabilities add up to a certain percentage (e.g., 0.9 means the model considers the top 90% most likely words, cutting off bizarre, low-probability words).  YouTube · Matt Williams +1
- **top_k** (Default: **40**): Limits the model to only choose from its top K most likely next words. A lower number (e.g., 10) keeps the output focused; a higher number (e.g., 50) allows for more diverse vocabulary.  Ollama +2
- **seed**: Sets a specific starting point for the model's random number generator. If you use the same seed, prompt, and parameters, you will get the exact same response every time.  YouTube · Matt Williams +2
- **repeat_penalty** (Default: **1.1**): Penalizes the model for repeating the same words or phrases too often. Higher numbers result in more diverse answers.  Ollama +1
- **repeat_last_n** (Default: **64**): Controls how far back the model looks in its own output to check for repetition.  Ollama +1

Here is a graphical representation of the effect of *temperature* vs *top_p*. Think of text generation as a funnel: **top_p cuts off** the unlikely choices to shape the pool, while **temperature reshapes the probabilities** of the remaining choices.




Figure 4: *top_p* vs temperature interaction

- Low *top_p* and temperature
 - Factual Q&As
 - Legal documents
 - Code generation
- High *top_p*, low temperature
 - Executive summaries
 - Sensitive content
- Low *top_p*, high temperature
 - Casual chatbots
 - Humorous tone
- High *top_p* and temperature
 - Stories
 - Poetry
 - Brainstorming

Figure 5: *top_p* vs temperature: when to use

Here is a video that explains very nicely the twin-concepts of *top_p* and *temperature*.




Figure 6: Excellent video on top_p vs temperature

Here are the four scenarios of interaction of these two parameters.

Four Common Interaction Scenarios

1. Low Temp + Low Top_p (The Strict Academic)

- **Settings:** temperature 0.2 , top_p 0.5
- **Interaction:** top_p eliminates everything except the top handful of highly predictable words. Then, temperature makes the highest-scoring word look even more dominant.
- **Result:** Highly repetitive, extremely safe, and deterministic text. This combo is perfect for code syntax, math, and data extraction where mistakes are not allowed.

2. High Temp + High Top_p (The Wild Brainstormer)

- **Settings:** temperature 1.2 , top_p 0.95
- **Interaction:** top_p leaves the door wide open, keeping almost every obscure and rare word in play. Then, temperature flattens the odds, allowing those rare, chaotic words to frequently win the selection lottery.
- **Result:** Highly creative, unpredictable, and poetic responses. However, it significantly increases the risk of hallucinating facts or generating grammatical gibberish.

3. High Temp + Low Top_p (The Controlled Novelty)

- **Settings:** temperature 1.2 , top_p 0.4
- **Interaction:** top_p strictly limits the selection to only the top 40% most logical words. However, because temperature is high, the model will aggressively cycle between those few valid options rather than just picking the #1 choice.
- **Result:** Vivid, lively writing that still adheres strictly to the topic without drifting off into nonsense.

Best Practices for Tuning

- **Adjust one at a time:** If you tweak both simultaneously, it is difficult to tell which parameter caused a shift in tone.
- **Keep one at default:** The industry standard recommendation is to leave top_p at its default value (0.9 or 0.95) and use temperature as your primary dial for creativity.
- **The Coding Override:** If you need absolute precision, set temperature to 0.0 . This forces the model to always pick the absolute highest probability token, rendering top_p completely irrelevant.

Custom Ollama Models

For Modelfile, please see this [GitHub link](#)
Opwebui for creating online modelfile [link is here](#)

Custom Ollama models are **personalized versions of existing AI models (like Llama 3 or Mistral) modified to behave in specific ways**. Created using a simple script called a **Modelfile**, they allow you to set custom system prompts, adjust behavior parameters, or load your own fine-tuned weights—all running locally on your hardware.  YouTube · Douglas Starnes +1



Key Ways to Customize Models

- **System Prompts & Personas:** Change the core identity of the model. For instance, you can turn a base model into a stern code reviewer, an expert copywriter, or a specific character.  YouTube · The Neural Maze +2
- **Parameters:** Adjust the settings for how the model generates responses. You can increase the `temperature` for more creative responses, adjust `num_ctx` to give it a larger memory/context window, or set `stop` sequences.  YouTube · The Neural Maze +2
- **Templates:** Dictate exactly how prompts are structured before being fed into the AI.  YouTube · The Neural Maze +1
- **Custom Weights (GGUF):** Import your own fine-tuned model weights or specialized models (like those from Hugging Face) to give the AI entirely new knowledge domains.  Ollama +1

How They Work: The Modelfile

A Modelfile is essentially a blueprint that tells Ollama how to construct your custom assistant, working much like a Dockerfile. Once written, you use the `ollama create` command to build it into an independent model that you can call upon anytime.  Ollama +3

Here is an example Modelfile:

```
FROM llama3.2
# sets the temperature to 1 [higher is more creative, lower is more coherent]
PARAMETER temperature 1
# sets the context window size to 4096. Control how many tokens the LLM can use
# as context to generate the next token
PARAMETER num_ctx 4096
# sets a custom system message to specify the behaviour of the chat assistant
SYSTEM You are Mario from super mario bros, acting as an assistant.
```

Creating custom ollama models using *Modelfile* are a pain. In most cases error comes. One way to avoid the error is to work from within docker. Also, Ollama *chat model* does not allow *temperature* to be set beyond 1.0. To do that, one can create ollama custom models with preset parameter values, as:

- Enter your container's terminal and open bash:**
`docker exec -it ollama /bin/bash`
(This also works → `docker exec -it ollama bash`)
- Create the Modelfile text directly on the container's hard drive:**
`echo -e "FROM llama3.2\nPARAMETER temperature 2.0" > Modelfile`
- Create custom model now:**
`ollama create custom-llama -f ./Modelfile`
- Modelfile is:**

```
FROM llama3.2
PARAMETER temperature 2.0
```

This works and custom model gets created (see figure below).


```

ashok@ashok:~$ docker exec -it ollama /bin/bash
root@ashok:/# echo -e "FROM llama3.2\nPARAMETER temperature 2.0" > Modelfile
root@ashok:/# cat Modelfile
FROM llama3.2
PARAMETER temperature 2.0
root@ashok:/# ollama create custom-llama -f ./Modelfile
gathering model components
using existing layer sha256:dde5aa3fc5ffc17176b5e8bdc82f587b24b2678c6c66101bf7da77af9f7ccdff
using existing layer sha256:966de95ca8a62200913e3f8bf84c8494536f1b94b49166851e76644e966396
using existing layer sha256:fcc5a6bec9daf9b561a68827b67ab6088e1dba9d1fa2a50d7bbcc8384e0a265d
using existing layer sha256:a70ff7e570d97baaf4e62ac6e6ad9975e04caa6d900d3742d37698494479e0cd
creating new layer sha256:dfe75703775a34b36542f3154b996fe054deef1032d02e66db089864d9941150
writing manifest
success

```

Figure 7: Ollama custom model creation from within docker

```

ashok@ashok:~$ ollama show custom-llama
Model
architecture      llama
parameters        3.2B
context length    131072
embedding length   3072
quantization       Q4_K_M

Capabilities
completion
tools

Parameters
stop               "<|start_header_id|>"
stop               "<|end_header_id|>"
stop               "<|eot_id|>"
temperature        2

License
LLAMA 3.2 COMMUNITY LICENSE AGREEMENT
Llama 3.2 Version Release Date: September 25, 2024

```

Figure 8: To see details of a model, use 'ollama show' command

Loop Engineering