

Technical concepts

What is an API?

An **API**, or **Application Programming Interface**, is a software intermediary that allows two different applications to interact and share data with one another. Think of it as a digital bridge or a translator that lets separate programs "speak" the same language without needing to know how each other's internal code is written. [Wikipedia +3](#)

The Restaurant Analogy

To understand how an API works, imagine you are a customer dining at a restaurant: [YouTube · Postman +1](#)

- **The Customer (The Client):** You sit at the table and want to order food. In technology, this is an app or website.
- **The Kitchen (The Server):** The kitchen prepares the food. This represents the backend database or system holding the data.
- **The Waiter (The API):** You don't walk into the kitchen to cook. Instead, you read the menu and give your order to the waiter. The waiter carries your request to the kitchen, retrieves the cooked food, and delivers it back to you. [GitHub +3](#)

How Web APIs Work

Most modern web communication relies on a structured **Request-and-Response** loop: [YouTube · Exponent +1](#)

1. **The Call:** Your application sends an **API request** over the internet using an HTTP protocol. [GitHub +1](#)
2. **The Endpoint:** The request travels to a specific URL path, known as an **endpoint**, which dictates what data is being asked for. [YouTube · CodeWithChris +1](#)
3. **The Method:** The request uses strict commands such as **GET** (to fetch data) or **POST** (to submit new data). [GitHub +1](#)
4. **The Security Check:** The system verifies the app's identity using a digital password called an **API key** to ensure the request is authorized. [YouTube · CodeWithChris +1](#)
5. **The Payload:** The server processes the task and sends back the data, typically formatted in a clean, machine-readable text format called **JSON**. [YouTube · Exponent +1](#)

What is a webhook?

The Doorbell Analogy

Think of a **traditional API** like a customer repeatedly calling a restaurant to ask, "Is my takeout order ready yet?". A **webhook** is like leaving your phone number with the restaurant so *they* can automatically text you the moment your food is ready.

Here, replace the 'customer' with a *web-server (say, App-A)*, 'restaurant' with another *web-server (say App-B)* and 'calling' with making an API request to 'App-B'; your 'phone number' with your 'URL or your API'. Thus, instead of 'App-A' sending an API request to 'App-B' for some information, 'App-B' makes an API request to 'App-A' informing him of an event.

Polling process vs. webhook




Figure 1: Instead of CRM, repeatedly sending a message to know if the customer order is dispatched, Warehouse, informs CRM as soon as customer order is dispatched.

API vs Webhook

APIs are pull-based: Your application (App-A) actively asks for data on-demand (you request it).

Webhooks are push-based: The source application (App-B) automatically sends the data to your application the moment something occurs

Webhook is an "Event-based HTTP callback trigger"—a long name. Name, "Webhook" is punchy, precise, and easily turned into a verb (e.g., "We can just *webhook* that data over"). Webhook uses standard web protocols (HTTP): POST/GET.

Ollama parameters

Ollama parameters are customizable settings that control how a large language model (LLM) generates text, balancing its creativity, predictability, memory, and formatting. They can be adjusted in the terminal, through an api/chat payload, or permanently saved in a **Modelfile**.

Key Generation Parameters

These parameters control the AI's "thinking" process and output style.

- **temperature** (Default: **0.8**): Controls the creativity and randomness of the output.
 - **Low** (e.g., **0.1** - **0.3**): Makes the output more factual, logical, and predictable. Best for coding, math, or data extraction.
 - **High** (e.g., **0.8** - **1.0**): Makes the output more creative, diverse, and unpredictable. Best for brainstorming or writing stories. Medium · Laurent Kubski +4
- **top_p** (Default: **0.9**): An alternative way to control randomness, often called "nucleus sampling". It tells the model to only consider the pool of words whose combined probabilities add up to a certain percentage (e.g., 0.9 means the model considers the top 90% most likely words, cutting off bizarre, low-probability words). YouTube · Matt Williams +1
- **top_k** (Default: **40**): Limits the model to only choose from its top K most likely next words. A lower number (e.g., **10**) keeps the output focused; a higher number (e.g., **50**) allows for more diverse vocabulary. Ollama +2
- **seed**: Sets a specific starting point for the model's random number generator. If you use the same seed, prompt, and parameters, you will get the exact same response every time. YouTube · Matt Williams +2
- **repeat_penalty** (Default: **1.1**): Penalizes the model for repeating the same words or phrases too often. Higher numbers result in more diverse answers. Ollama +1
- **repeat_last_n** (Default: **64**): Controls how far back the model looks in its own output to check for repetition. Ollama +1

Here is a graphical representation of the effect of temperature vs top_p. Think of text generation as a funnel: **top_p cuts off** the unlikely choices to shape the pool, while **temperature reshapes the probabilities** of the remaining choices.




Figure 2: top_p vs temperature interaction

- Low top_p and temperature
 - Factual Q&As
 - Legal documents
 - Code generation
- High top_p, low temperature
 - Executive summaries
 - Sensitive content
- Low top_p, high temperature
 - Casual chatbots
 - Humorous tone
- High top_p and temperature
 - Stories
 - Poetry
 - Brainstorming

Figure 3: top_p vs temperature: when to use



Figure 4: Excellent video on top_p vs temperature

Four Common Interaction Scenarios

1. Low Temp + Low Top_p (The Strict Academic)

- **Settings:** `temperature 0.2`, `top_p 0.5`
- **Interaction:** `top_p` eliminates everything except the top handful of highly predictable words. Then, `temperature` makes the highest-scoring word look even more dominant.
- **Result:** Highly repetitive, extremely safe, and deterministic text. This combo is perfect for code syntax, math, and data extraction where mistakes are not allowed.

2. High Temp + High Top_p (The Wild Brainstormer)

- **Settings:** `temperature 1.2`, `top_p 0.95`
- **Interaction:** `top_p` leaves the door wide open, keeping almost every obscure and rare word in play. Then, `temperature` flattens the odds, allowing those rare, chaotic words to frequently win the selection lottery.
- **Result:** Highly creative, unpredictable, and poetic responses. However, it significantly increases the risk of hallucinating facts or generating grammatical gibberish.

3. High Temp + Low Top_p (The Controlled Novelty)

- **Settings:** `temperature 1.2`, `top_p 0.4`
- **Interaction:** `top_p` strictly limits the selection to only the top 40% most logical words. However, because `temperature` is high, the model will aggressively cycle between those few valid options rather than just picking the #1 choice.
- **Result:** Vivid, lively writing that still adheres strictly to the topic without drifting off into nonsense.

Best Practices for Tuning

- **Adjust one at a time:** If you tweak both simultaneously, it is difficult to tell which parameter caused a shift in tone.
- **Keep one at default:** The industry standard recommendation is to leave `top_p` at its default value (`0.9` or `0.95`) and use `temperature` as your primary dial for creativity.
- **The Coding Override:** If you need absolute precision, set `temperature` to `0.0`. This forces the model to always pick the absolute highest probability token, rendering `top_p` completely irrelevant.